



Handling Asynchronicity

COMPUTE

| STORE


| ANALYZE



Exploiting overlap

- **If individual kernels are performing well**
 - Can you start overlapping different computational tasks?
 - kernels on the GPU
 - data transfers to/from the GPU
 - maybe even separate computation on the CPU
- **You can do this using the async features**
 - **OpenACC**: stream handles; very similar to streams in **CUDA**
 - **OpenMP**: data dependencies; same as OpenMP tasks

Asynchronicity with Nvidia GPUs



- **accelerator operations are launched from the CPU**
- **they then execute asynchronously**
 - control returns immediately to the host
 - host must then wait or test for completion
 - automatically (e.g. handled by compiler)
 - manually (e.g. via directives or API calls)
 - applies to:
 - computational kernels
 - data transfers: **update** directives
 - plus some other directives



Automatic synchronisation

- **By default, the compiler will handle synchronisation:**
 - may be conservative, and wait for every operation to complete; or
 - may be smarter, and reduce number of sync. points (consistent with correctness)
- **CCE-specific options:**
 - control behaviour with compiler option `-hacc_model=auto_async_*`
 - `auto_async_none` :
 - waits for every operation to separately complete
 - useful for debugging and generating profiles
 - but it may skew performance
 - `auto_async_kernels` :
 - may allow some computational kernels to overlap
 - the default behaviour
 - `auto_async_all` :
 - may allow kernels and data transfers to overlap
 - try this as a performance-tuning option

A stream of tasks in OpenMP 4.0 (and 4.5)

- Can simply add **nowait** clause to kernels, updates
 - Over-rides automatic synchronisation
- **Operations:**
 - guaranteed to execute sequentially in the order they were launched
 - not guaranteed to execute immediately after each other; there could be delays
 - this is known as a "stream" of tasks
- **Synchronisation**
 - **taskwait** or **barrier** directives ensure all asynchronous operations have completed
- If you've used **CUDA** streams, these concepts should be very familiar

Streams of tasks in OpenMP 4.5



- **Asynchronicity handled in the same way as tasks (4.0)**
 - Each kernel or set of data transfers is a task
- **depend clause gives much greater control than nowait**
 - Over-rides automatic synchronisation
 - Allows user to specify dependencies between tasks
 - These control:
 - order of execution
 - whether tasks can overlap or not



Dependency-types

- Dependencies are expressed as data dependencies
 - `depend(dependency-type:variable)`
- Dependency types are:
 - `in`:
 - `in:a` task cannot start until all tasks with `out:a` or `inout:a` have finished
 - `out`:
 - `out:a` task must complete before any `in:a` or `inout:a` tasks can start
 - `inout`: combination of `in` and `out` behaviour for same variable



Dependency variables

- **in/out/inout** dependency types
 - Simply used to define a dependency tree
 - Do not imply any data movement to/from the accelerator
 - Do not express how the variables is used in the task
 - **in:a** does not mean **a** is used in a read-only fashion in the task
 - **a** could be used in a write-only or read-write fashion, or
 - **a** might not be referenced at all in the task

OpenMP depend first example

- **a simple pipeline:**

- processing an array, slice by slice
 - assume can be processed independently
- each slice requires 3 tasks:
 - copy data to GPU,
 - process on GPU,
 - bring back to CPU
- which must execute sequentially
- but we can overlap different slices

- **Expect to see overlap of three streams at once**

- one sending to the device; one processing the slice; one sending to host

```
REAL :: a(Nvec,Nchunk),b(Nvec,Nchunk)

!$omp target data map(alloc:a,b)
DO j = 1,Nchunks
!$omp target update to(a(:,j)) &
    depend(out:a(:,j))

!$omp target teams distribute &
    depend(inout:a(:,j))
    DO i = 1,Nvec
        b(i,j) = <function of a(:,j)>
    ENDDO
!$omp end target teams distribute

!$omp target update from(b(:,j)) &
    depend(in:a(:,j))

ENDDO
!$omp taskwait
!$omp end target data
```

High-level example with OpenMP 4.5



```
!$omp target depend(out:x)
<Kernel A>
!$omp target depend(inout:x)
<Kernel B>

!$omp target depend(out:y)
<Kernel C>


!$omp target depend(inout:x) &
    depend(inout:y)
<Kernel D>

!$omp target depend(in:x)
<Kernel E>

!$omp target depend(inout:y)
<Kernel F>

!$omp target depend(in:y)
<Kernel G>

!$omp taskwait ! all completed
```



COMPUTE

STORE

ANALYZE



Synchronising subsets of tasks

- Dependencies ensure that tasks execute in order
- Can also globally synchronise
 - taskwait and barrier directives
- May also want to add CPU task into stream, e.g.
 - kernel to pack buffer on accelerator
 - update to move buffer back to host



A stream of tasks in OpenACC

- Add **async** clause to make operation asynchronous
 - kernel or data transfer
- **Operations:**
 - guaranteed to execute sequentially in the order they were launched
 - not guaranteed to execute immediately after each other; there could be delays
 - this is known as a "stream" of tasks
- **Synchronisation**
 - **wait** directive ensures all asynchronous operations have completed
 - can use an API call to do same thing (or to test for completion)
- If you've used **CUDA** streams, these concepts should be very familiar



Multiple streams of tasks in OpenACC

- The **async** clause can take a handle
 - handle is an argument that is a positive- or zero-valued integer
 - each handle value is a separate stream of tasks
- **Tasks within a given stream**
 - guaranteed to execute sequentially in order they were issued
- **Tasks in different streams**
 - may overlap or serialise, as the hardware and runtime allows
 - operations in different streams should be independent or we have a race condition
- **Synchronisation**
 - **wait** directive ensures all streams of tasks have completed
 - **wait(handle)** directive ensures just the specified stream of tasks has completed
 - can use an API call to do same thing (or to test for completion)
- If you've used **CUDA** streams, these concepts should be very familiar

OpenACC async first example



- **a simple pipeline:**

- processing an array, slice by slice
- assume slices can be processed independently
- each slice requires 3 tasks:
 - copy data to GPU,
 - process on GPU,
 - bring back to CPU
- which must execute sequentially
- but we can overlap different slices

- Use a different stream for each slice
 - use slice number as stream handle
 - don't worry if number gets too large

```
REAL :: a(Nvec,Nchunk),b(Nvec,Nchunk)

!$acc data create(a,b)
DO j = 1,Nchunks
!$acc update device(a(:,j)) async(j)

!$acc parallel loop async(j)
  DO i = 1,Nvec
    b(i,j) = <function of a(:,j)>
  ENDDO

!$acc update host(b(:,j)) async(j)

ENDDO
!$acc wait
!$acc end data
```

- **Expect to see overlap of three streams at once**

- one sending to the device; one processing the slice; one sending to host


OpenACC async results

- **Execution times:**

- CPU: 3.76s
- OpenACC: 1.10s
- OpenACC, async: 0.34s

- **NVIDIA Visual profiler**

- only 7 of the 16 streams shown
- collapsed view at bottom



High-level example with OpenACC

```

!$acc parallel loop async(stream1)
<Kernel A>
!$acc parallel loop async(stream1)
<Kernel B>

!$acc parallel loop async(stream2)
<Kernel C>


!$acc parallel loop async(stream3) &
!$acc      wait(stream1,stream2)
<Kernel D>

!$acc parallel loop async(stream4) &
!$acc      wait(stream3)
<Kernel E>

!$acc parallel loop async(stream5) &
!$acc      wait(stream3)
<Kernel F>

!$acc parallel loop async(stream5)
<Kernel G>

!$acc wait ! ensures all completed
    
```



What's missing in OpenACC and OpenMP?



- Deep copy is the elephant in the room
- Hierarchical data structures with pointers
 - C++ objects; C structs; Fortran derived types
 - On CPU, pointers point to CPU memory addresses
 - When "map" to GPU, pointers still point to CPU addresses
 - main impact: at best, inefficient; usually, just broken
 - User needs to explicitly remap all the pointers
 - this is not a very satisfactory solution
- CCE **-hacc_model=deep_copy** helps for Fortran only
- Proper solution deferred to OpenMP 5.0 (due Nov.17)
 - and OpenACC 3.0 (due ?)

