



# DIRECTIVE BASED GPU PROGRAMMING WORKSHOP @ CSCS

Markus Wetzstein, 31.05.2017



# OPENACC ADOPTION





# OPENACC ADOPTION

## NEW PLATFORMS



Sunway TaihuLight  
#1 Top 500, Nov. 2016

## GROWING COMMUNITY



- 6,000+ enabled developers
- 4,500+ course registrants
- 250+ Hackathon attendees
- 150+ User Group members

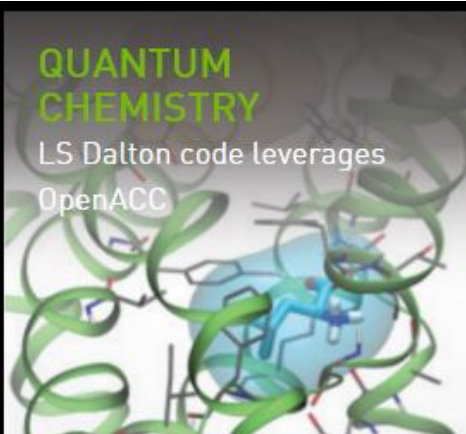
## PORTING SUCCESS

- **Gaussian 16** ported to Tesla with OpenACC
- Five of 13 **ORNL CAAR** codes using OpenACC
- **ANSYS Fluent R18** production release is GPU accelerated using OpenACC

# OPENACC SUCCESS STORIES

## QUANTUM CHEMISTRY

LS Dalton code leverages OpenACC



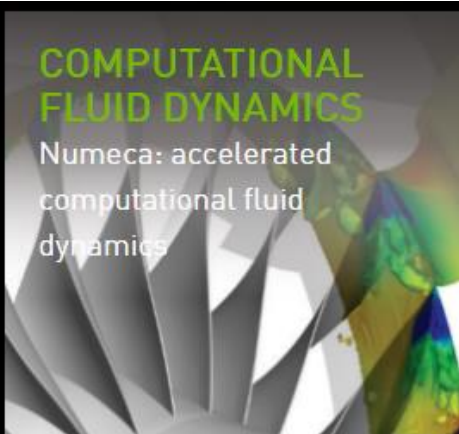
## MEDICAL IMAGING

Accelerated MRI reconstruction model



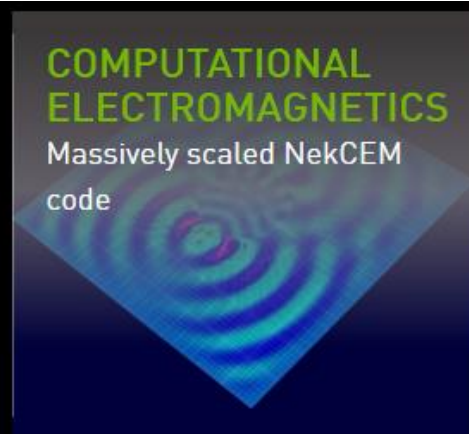
## COMPUTATIONAL FLUID DYNAMICS

Numeca: accelerated computational fluid dynamics



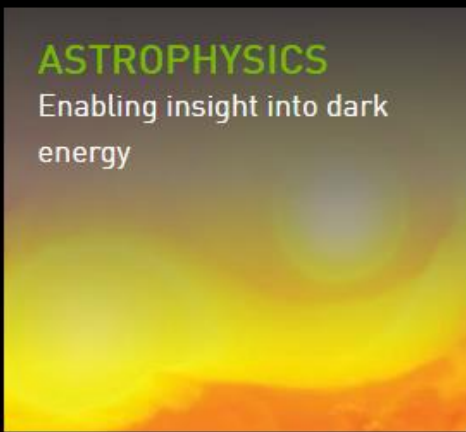
## COMPUTATIONAL ELECTROMAGNETICS

Massively scaled NekCEM code



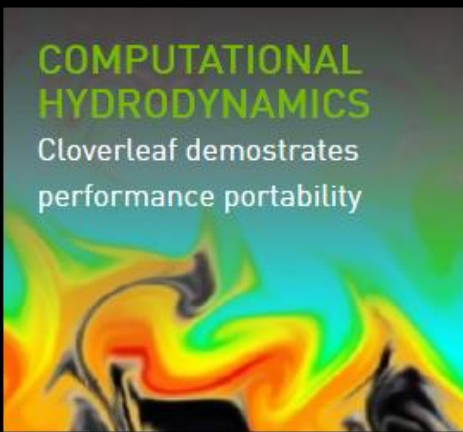
## ASTROPHYSICS

Enabling insight into dark energy



## COMPUTATIONAL HYDRODYNAMICS

Cloverleaf demonstrates performance portability



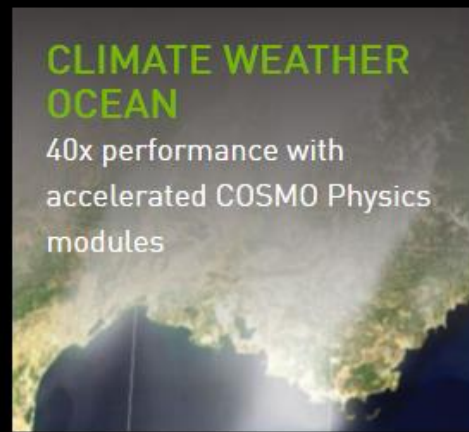
## COMPUTATIONAL FLUID DYNAMICS

Fully implicit 3D CFD solver, 4x in 5 days



## CLIMATE WEATHER OCEAN

40x performance with accelerated COSMO Physics modules



[developer.nvidia.com/openacc/success-stories](https://developer.nvidia.com/openacc/success-stories)

# THE NVIDIA HPC SDK

PGI Fortran, C & C++ Compilers

Optimizing, SIMD Vectorizing, OpenMP

Accelerated Computing Features

OpenACC Directives

CUDA Fortran

Multi-Platform Solution

Multicore x86-64 and OpenPOWER CPUs,  
NVIDIA Tesla GPUs

Supported on Linux, macOS, Windows

MPI/OpenMP/OpenACC Tools

Debugger

Performance Profiler

Interoperable with DDT, TotalView

# PGI<sup>®</sup>

The Compilers & Tools  
for Supercomputing



# OPENACC FOR EVERYONE

## New PGI Community Edition Now Available

|  | <b>FREE</b><br><b>PGI</b><br>Community<br>EDITION | <b>PGI</b><br>Professional<br>EDITION | <b>PGI</b><br>Enterprise<br>EDITION |
|--|---|---------------------------------------|-------------------------------------|
| <b>PROGRAMMING MODELS</b><br>OpenACC, CUDA Fortran, OpenMP,<br>C/C++/Fortran Compilers and Tools | ✓   | ✓                                     | ✓                                   |
| <b>PLATFORMS</b><br>X86, OpenPOWER, NVIDIA GPU   | ✓   | ✓                                     | ✓                                   |
| <b>UPDATES</b>   | 1-2 times a year                                  | 6-9 times a year                      | 6-9 times a year                    |
| <b>SUPPORT</b>   | User Forums                                       | PGI Support                           | PGI Premier Services                |
| <b>LICENSE</b>   | Annual  | Perpetual                             | Volume/Site                         |



# USING THE PGI COMPILER SUITE



# INVOKING THE COMPILER

On Cray machines like Piz Daint @ CSCS:

```
module load PrgEnv-pgi
module avail pgi
module load pgi/...
module load craype-accel-nvidiaXX
```

The wrappers will then point to the PGI compilers: `ftn`, `cc`, `CC`

Without Cray wrappers: `pgfortran`, `pgcc`, `pgc++`

Useful basic options: `-fast` for optimizations, `-Minfo` for compiler feedback

If you want to see the possible arguments to a given option and check what the default is: `pgf90 -help [some other option]`



# BUILDING ACCELERATOR PROGRAMS

```
pgfortran -acc a.f90
```

```
pgcc -acc a.c
```

```
pgc++ -acc a.cpp
```

Other options:

```
-ta=tesla[:cc2x|cc3x|cc50|cc60]
```

```
-ta=tesla[:cuda7.0|cuda7.5|cuda8.0]
```

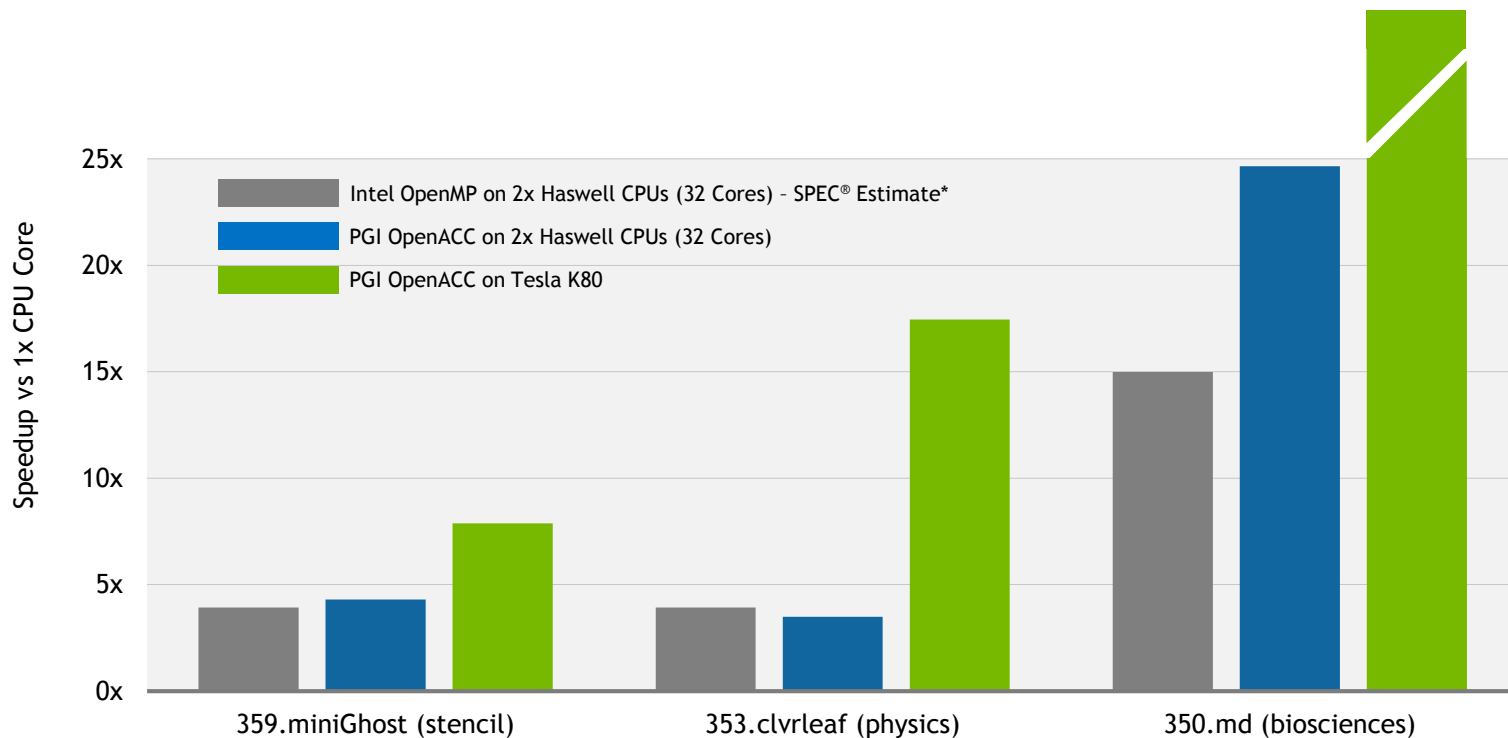
```
-ta=tesla,host [default with -acc]
```

```
-ta=multicore
```

RECOMMENDED: Enable compiler feedback with **-Minfo** or **-Minfo=accel**

# PGI OPENACC FOR MULTI-CORE X86 CPUS

OpenACC Performance Portability Across Multicore x86 CPUs and GPUs



Supermicro SYS-2028GR-TRT, Intel Xeon E5-2698 v3, 32 cores, NVIDIA Tesla K80, 256GB of System Memory

PGI 15.7 Beta OpenACC Multicore and K80 results from SPEC ACCEL™ measured June 2015.

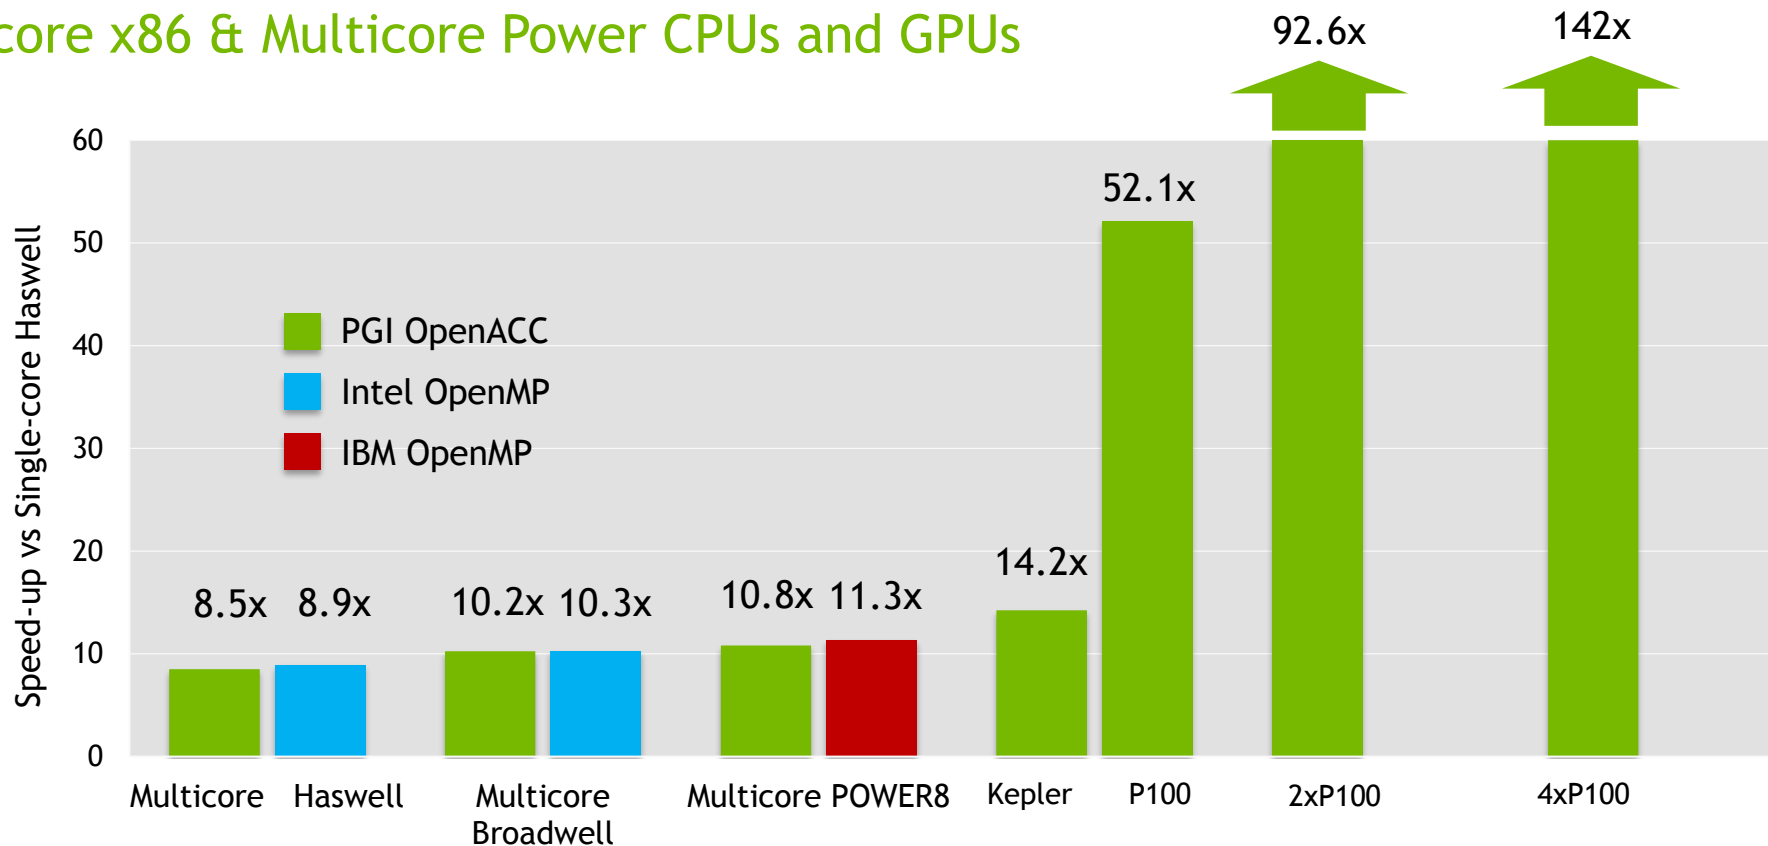
\* Intel 15.0.90 OpenMP results use Cloverleaf reference application and SPEC OMP®2012 using workloads from SPEC ACCEL.

SPEC® and the benchmark names SPEC ACCEL™ and SPEC OMP® are registered trademarks of the Standard Performance Evaluation Corporation.

More info: SPEC ACCEL [www.spec.org/accel](http://www.spec.org/accel), OMP2012 [www.spec.org/omp2012](http://www.spec.org/omp2012), CloverLeaf OpenMP\_uk-mac [github.io/CloverLeaf/](https://github.io/CloverLeaf/), miniGhost ref 1.0.1 <https://mantevo.org/download/>

# OPENACC PERFORMANCE PORTABILITY - CLOVERLEAF 1.3

Multicore x86 & Multicore Power CPUs and GPUs



AWE Hydrodynamics CloverLeaf mini-App, bm32 data set



# PGI\_ACC\_TIME ENVIRONMENT VARIABLE

Accelerator Kernel Timing data

/proj/scratch/mwolfe/test/openacc/src/smooth4.c

smooth\_acc NVIDIA devicenum=0

time(us): 317

12: data region reached 5 times

12: data copyin reached 10 times

device time(us): total=121 max=19 min=11 avg=12

23: data copyout reached 5 times

device time(us): total=63 max=14 min=12 avg=12

14: compute region reached 5 times

17: kernel launched 5 times

grid: [1x98] block: [128]

device time(us): total=133 max=90 min=9 avg=26

elapsed time(us): total=176 max=99 min=17 avg=35

# PGI\_ACC\_NOTIFY BIT MASK

## 1 - launch

```
launch CUDA kernel  file=smooth4.c function=smooth_acc line=17 device=0 num_gangs=98  
num_workers=1 vector_length=128 grid=1x98 block=128
```

## 2 - data upload/download

```
upload CUDA data  file=smooth4.c function=smooth_acc line=12 device=0 variable=a  
bytes=40000  
download CUDA data  file=smooth4.c function=smooth_acc line=23 device=0 variable=a  
bytes=40000
```

## 4 - wait (explicit or implicit) for device

```
Implicit wait  file=smooth4.c function=smooth_acc line=17 device=0  
Implicit wait  file=smooth4.c function=smooth_acc line=23 device=0
```

## 8 - data/compute region enter/leave

```
Enter data region file=smooth4.c function=smooth_acc line=12 device=0  
Enter compute region file=smooth4.c function=smooth_acc line=14 device=0  
Leave compute region file=smooth4.c function=smooth_acc line=17 device=0
```

## 16 - data create/allocate/delete/free

```
create CUDA data  bytes=40000 file=smooth4.c function=smooth_acc line=12 device=0  
alloc  CUDA data  bytes=40000 file=smooth4.c function=smooth_acc line=12 device=0  
delete CUDA data  bytes=40448 file=smooth4.c function=smooth_acc line=23 device=0
```

# PGPROF COMMAND LINE PROFILER

`pgprof ./exe`

Report kernel and transfer times directly

Collect profiles for PGPROF GUI

`%> pgprof -o profile.out ./exe`

Collect for MPI processes

`%> mpirun -np 2 pgprof -o profile.%p.out ./exe`

Collect profiles for complex process hierarchies

`--profile-child-processes, --profile-all-processes`

Collect key events and metrics

`%> pgprof -events all flops_sp ./exe`

Trace stream usage

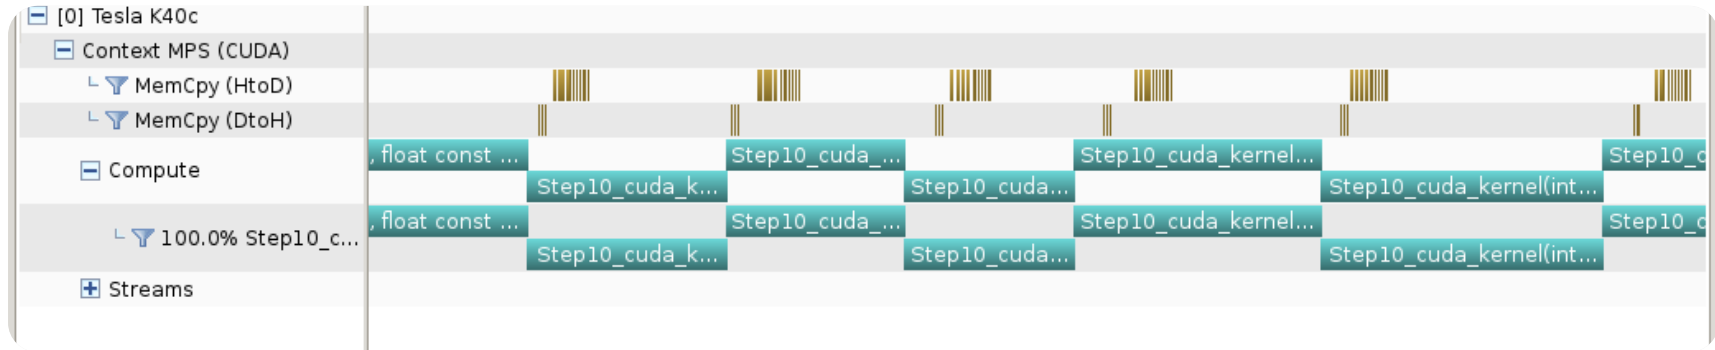
`%> pgprof -print-gpu-trace ./exe`

Full listing of options see: `pgprof --help`



# PGPROF VISUAL PROFILER

## Timeline



## Guided System

**1. CUDA Application Analysis**

**2. Performance-Critical Kernels**

**3. Compute, Bandwidth, or Latency Bound**

The first step in analyzing an individual kernel is to determine if the performance of the kernel is bounded by computation, memory bandwidth, or instruction/memory latency. The results at right indicate that the performance of kernel "Step10\_cuda\_kernel" is most likely limited by compute.

[Perform Compute Analysis](#)

The most likely bottleneck to performance for this kernel is compute so you should first perform compute analysis to determine how it is limiting performance.

[Perform Latency Analysis](#)

[Perform Memory Bandwidth Analysis](#)

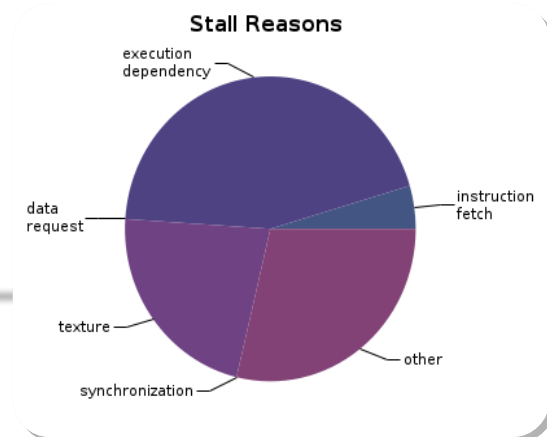
Instruction and memory latency and memory bandwidth are likely not the primary performance bottlenecks for this kernel, but you may still want to perform those analyses.

[Rerun Analysis](#)

If you modify the kernel you need to rerun your application to update this analysis.

## Analysis

| Shared Memory  |         |              |
|--|---------|--------------|
| Local Loads  | 0       | 0 B/s        |
| Local Stores   | 0       | 0 B/s        |
| Shared Loads   | 0       | 0 B/s        |
| Shared Stores  | 0       | 0 B/s        |
| Global Loads   | 0       | 0 B/s        |
| Global Stores  | 0       | 0 B/s        |
| L1/Shared Total  | 0       | 0 B/s        |
| L2 Cache   |         |              |
| Reads  | 6339426 | 236.738 GB/s |
| Writes   | 31414   | 1.173 GB/s   |
| Total  | 6370840 | 237.912 GB/s |
| Texture Cache  |         |              |
| Reads  | 6450496 | 240.886 GB/s |
| Device Memory  |         |              |
| Reads  | 1562634 | 58.355 GB/s  |
| Writes   | 7504    | 280.228 MB/s |
| Total  | 1570138 | 58.635 GB/s  |
| System Memory [ PCIe configuration: Gen3 x16, @ 8 Gbit/s ] |         |              |
| Reads  | 0       | 0 B/s        |
| Writes   | 4       | 149.375 kB/s |
| Total  | 4       | 149.375 kB/s |





# OPENACC DIRECTIVES: ROUTINE DECLARE ATOMIC

# OPENACC **routine** DIRECTIVE

**routine:** Compile the following function for the device (allows a function call in device code)

Clauses: gang, worker, vector, seq

```
#pragma acc routine seq
void fun(...) {
    for(int i=0;i<N;i++)
        ...
}
```

```
#pragma acc routine vector
void fun(...) {
    #pragma acc loop vector
    for(int i=0;i<N;i++)
        ...
}
```



# OPENACC **routine** DIRECTIVE

**routine:** Compile the following function for the device (allows a function call in device code)

Clauses: gang, worker, vector, seq

```
#pragma acc parallel loop gang \
    vector_length(VL)
for(int i=0;i<N;i++)
    fun_vec(...);
}
```

```
#pragma acc routine vector
void fun_vec(...) {
    #pragma acc loop vector
    for(int i=0;i<N;i++)
        fun_seq(...);
}
```

```
#pragma acc routine seq
void fun_seq(...) {

    for(int i=0;i<N;i++)
        ...
}
```

# OPENACC **routine**: FORTRAN

```
subroutine foo(v, i, n) {  
  use ...  
  !$acc routine vector  
  real :: v(:, :)  
  integer, value :: i, n  
  !$acc loop vector  
  do j=1,n  
    v(i,j) = 1.0/(i*j)  
  enddo  
end subroutine  
  
$!acc parallel loop  
do i=1,n  
  call foo(v,i,n)  
enddo  
$!acc end parallel loop
```

The **routine** directive may appear in a Fortran function or subroutine definition, or in an interface block.

Nested acc routines require the routine directive within each nested routine.

The save attribute is not supported.

Note: Fortran, by default, passes all arguments by reference. Passing scalars by value will improve performance of GPU code.

# SEQ ROUTINE AUTO-GENERATION

In a **C++ program**, many functions, certainly class member functions, appear as source code in header files included in the program. Oftentimes, the functions are defined in system headers or other application packages, and modifying those headers is either unwise or impossible.

The PGC++ compiler will take note of functions called in compute regions and implicitly add the pragma **acc routine seq** if there is no explicit **routine** directive.



# GLOBAL DATA

```
float a[1000000];
```

```
extern void matvec(...);
```

```
...
```

```
for (i = 0; i < m; ++i) {  
    matvec(v, x, i, n);  
}
```

```
extern float a[];
```

```
void matvec (float* v, float* x, int  
I, int n) {  
    for (int j = 0; j < n; ++j)  
        x[i] += a[i*n+j] * v[j];  
}
```

# DECLARE CREATE

```
float a[1000000];  
#pragma acc declare create(a)  
  
#pragma acc routine worker  
extern void matvec(...);  
...  
  
#pragma acc parallel loop  
for (i = 0; i < m; ++i) {  
    matvec(v, x, i, n);  
}
```

```
extern float a[];  
#pragma acc declare create(a)  
  
#pragma acc routine worker  
void matvec (float* v, float* x, int  
I, int n) {  
    #pragma acc loop worker  
    for (int j = 0; n < n; ++j)  
        x[i] += a[i*n+j] * v[j];  
}
```

# GLOBAL DATA

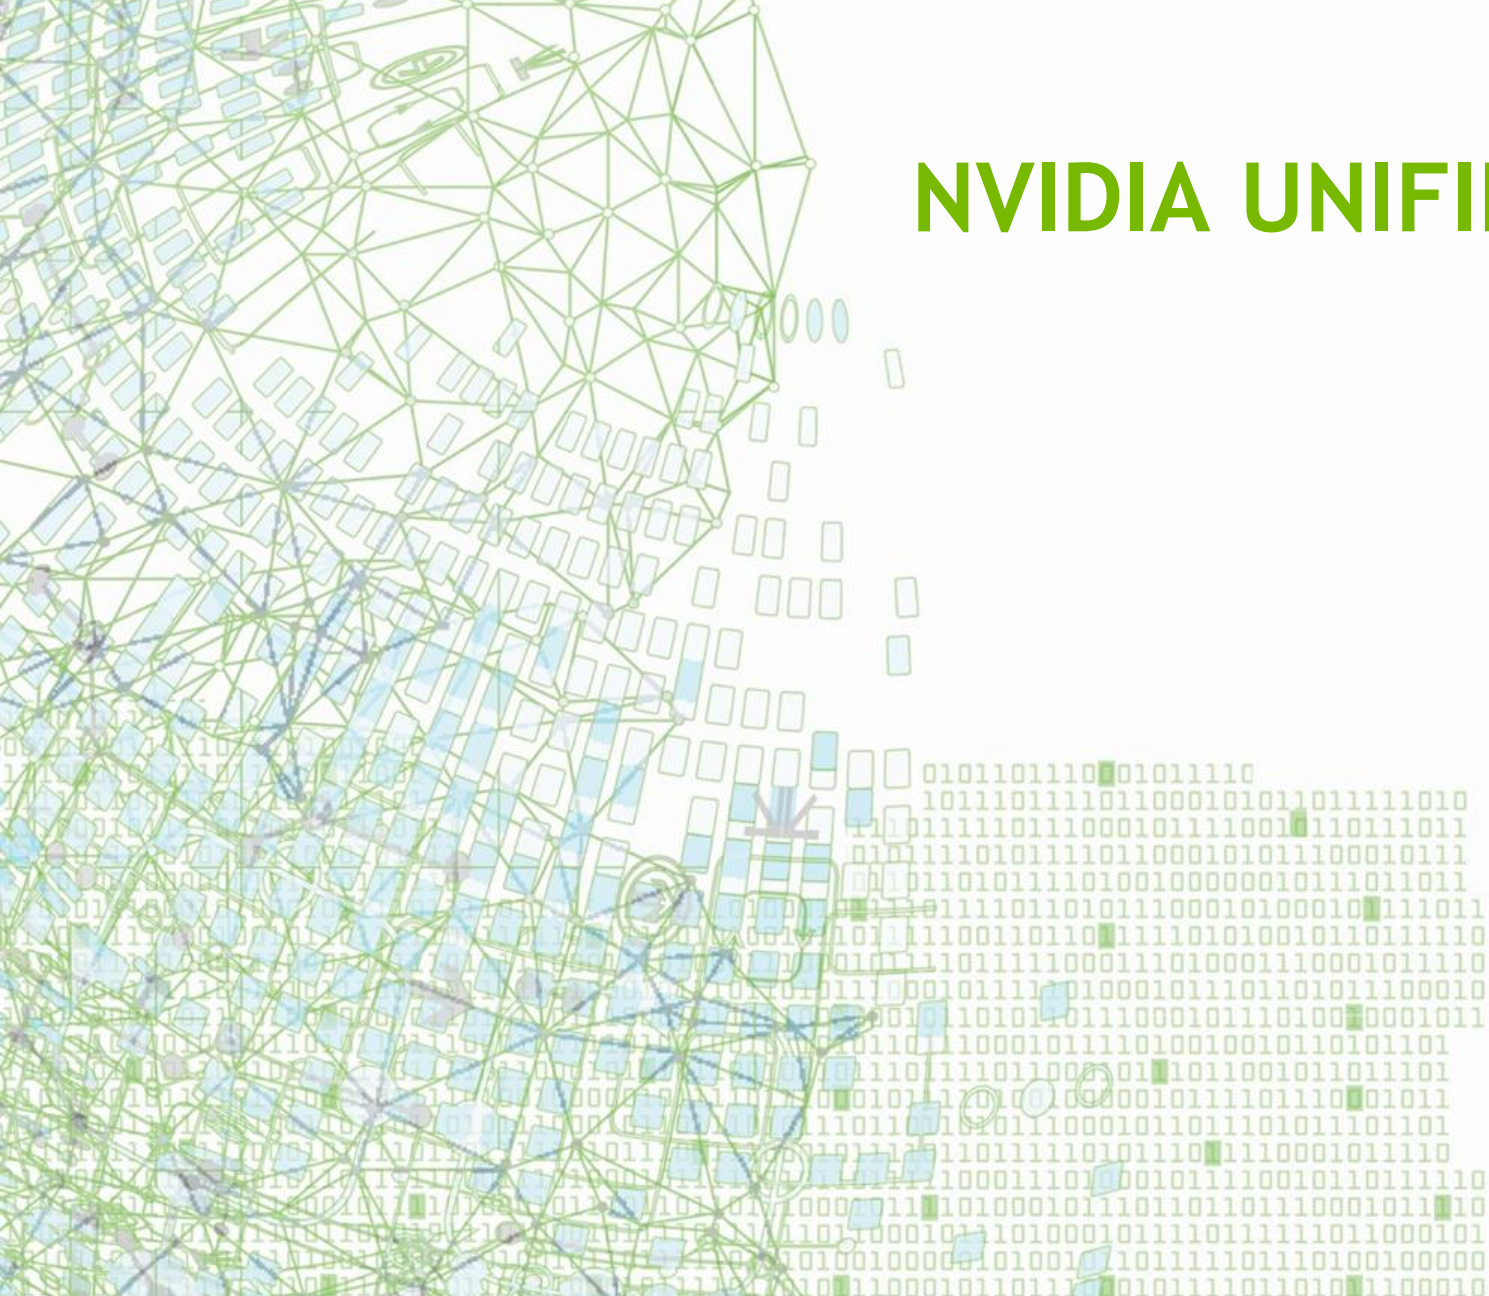
```
Fortran:      module m
               real, allocatable :: x(:)
               !$acc declare create(x)
               ! x dynamically allocated on host+device
               end module
```

Summary **declare** directive:

Used to generate a device declaration for global data.

There are more complicated use cases and additional clauses to **declare**, but **declare create** is most likely what you will need for most scenarios.

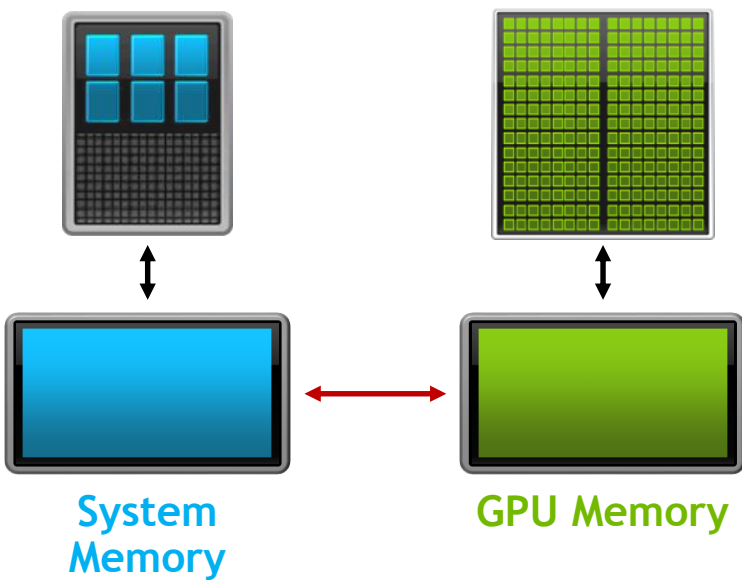
# NVIDIA UNIFIED MEMORY



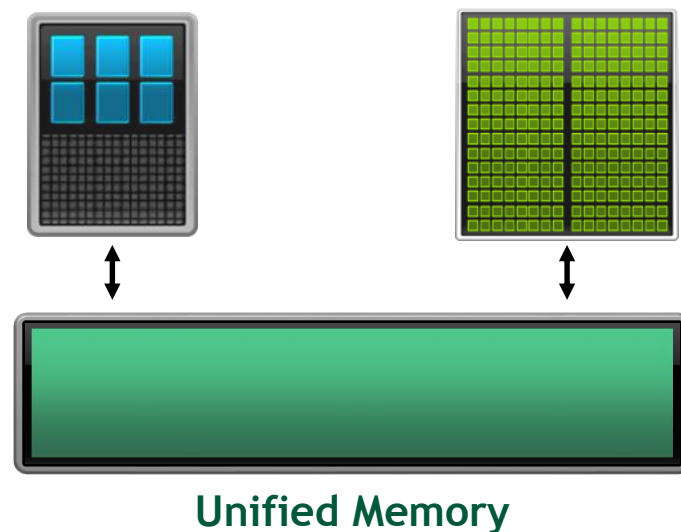
# UNIFIED MEMORY

Dramatically Lower Developer Effort

Developer View Today



Developer View With Unified Memory





# TESLA GPU PROGRAMMING IN 3 STEPS

## PARALLELIZE

Parallelize with OpenACC  
for multicore CPUs

```
% pgc++ -ta=multicore ...
```

```
while ( error > tol && ...  
    error = 0.0;  
#pragma acc parallel loop ...  
    for( int j = 1; ...  
#pragma acc loop  
    for( int i = 1; ...  
        ...  
    }  
...  
...
```

## OFFLOAD

Port to Tesla using OpenACC  
with CUDA Unified Memory

## OPTIMIZE

Optimize and overlap data  
movement using OpenACC  
data directives

# TESLA GPU PROGRAMMING IN 3 STEPS

## PARALLELIZE

Parallelize with OpenACC  
for multicore CPUs

```
% pgc++ -ta=multicore ...
```

```
while ( error > tol && ...  
    error = 0.0;  
#pragma acc parallel loop ...  
    for( int j = 1; ...  
#pragma acc loop  
    for( int i = 1; ...  
        ...  
    }  
...  
...
```

## OFFLOAD

Port to Tesla using OpenACC  
with CUDA Unified Memory

```
% pgc++ -ta=tesla:managed ...
```

```
while ( error > tol && ...  
    error = 0.0;  
#pragma acc parallel loop ...  
    for( int j = 1; ...  
#pragma acc loop  
    for( int i = 1; ...  
        ...  
    }  
...  
...
```

## OPTIMIZE

Optimize and overlap data  
movement using OpenACC  
data directives

# TESLA GPU PROGRAMMING IN 3 STEPS

## PARALLELIZE

Parallelize with OpenACC  
for multicore CPUs

```
% pgc++ -ta=multicore ...
```

```
while ( error > tol && ...  
    error = 0.0;  
#pragma acc parallel loop ...  
    for( int j = 1; ...  
#pragma acc loop  
        for( int i = 1; ...  
            ...  
        }  
...  
...
```

## OFFLOAD

Port to Tesla using OpenACC  
with CUDA Unified Memory

```
% pgc++ -ta=tesla:managed ...
```

```
while ( error > tol && ...  
    error = 0.0;  
#pragma acc parallel loop ...  
    for( int j = 1; ...  
#pragma acc loop  
        for( int i = 1; ...  
            ...  
        }  
...  
...
```

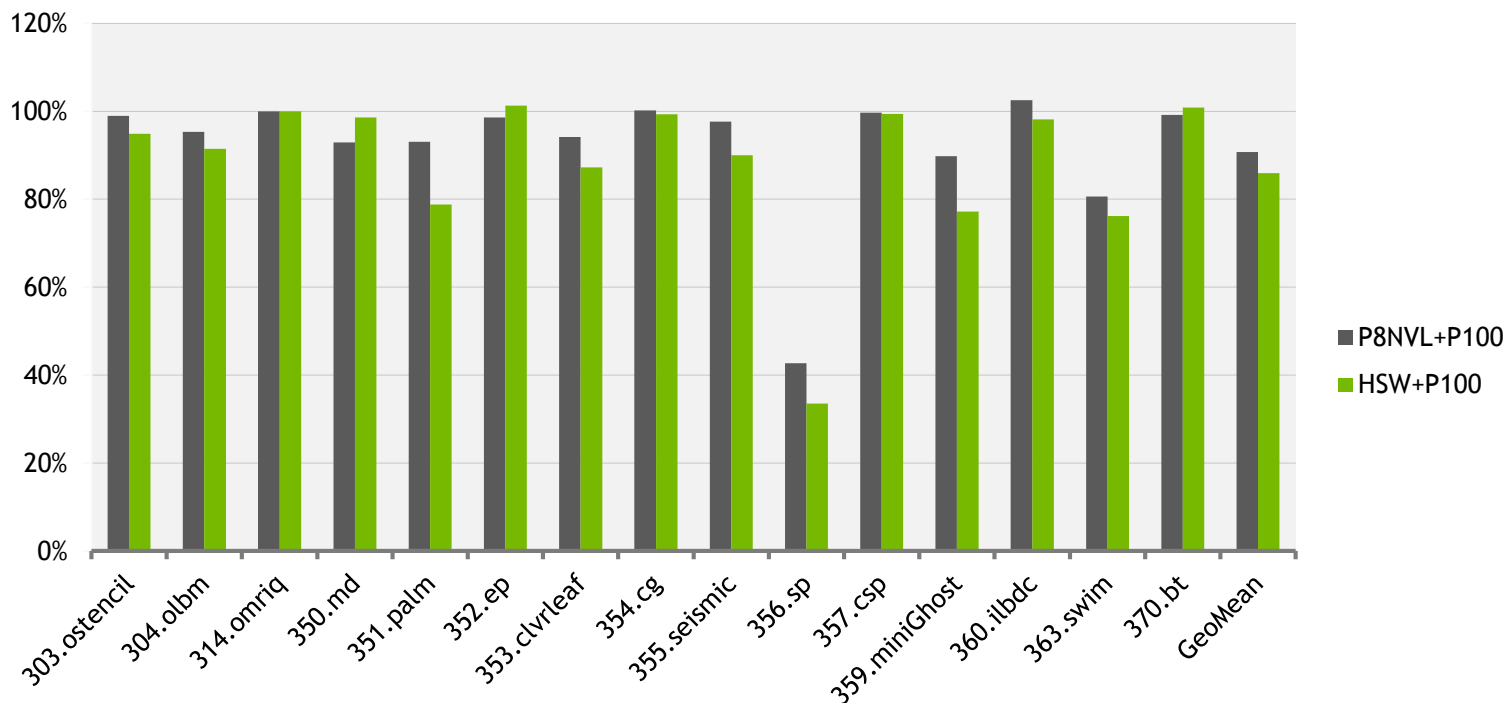
## OPTIMIZE

Optimize and overlap data  
movement using OpenACC  
data directives

```
#pragma acc data create ...  
while ( error > tol && ...  
    error = 0.0;  
#pragma acc parallel loop ...  
    for( int j = 1; ...  
#pragma acc loop  
        for( int i = 1; ...  
            ...  
        }  
...  
...
```

# CUDA UNIFIED MEMORY ON NVLINK VS PCIE

P100 Paging Engine Moves All Dynamically Allocated Data



100% = Directive-based  
Data Movement

OpenACC w/CUDA 8.0 UM using  
-ta=tesla:managed

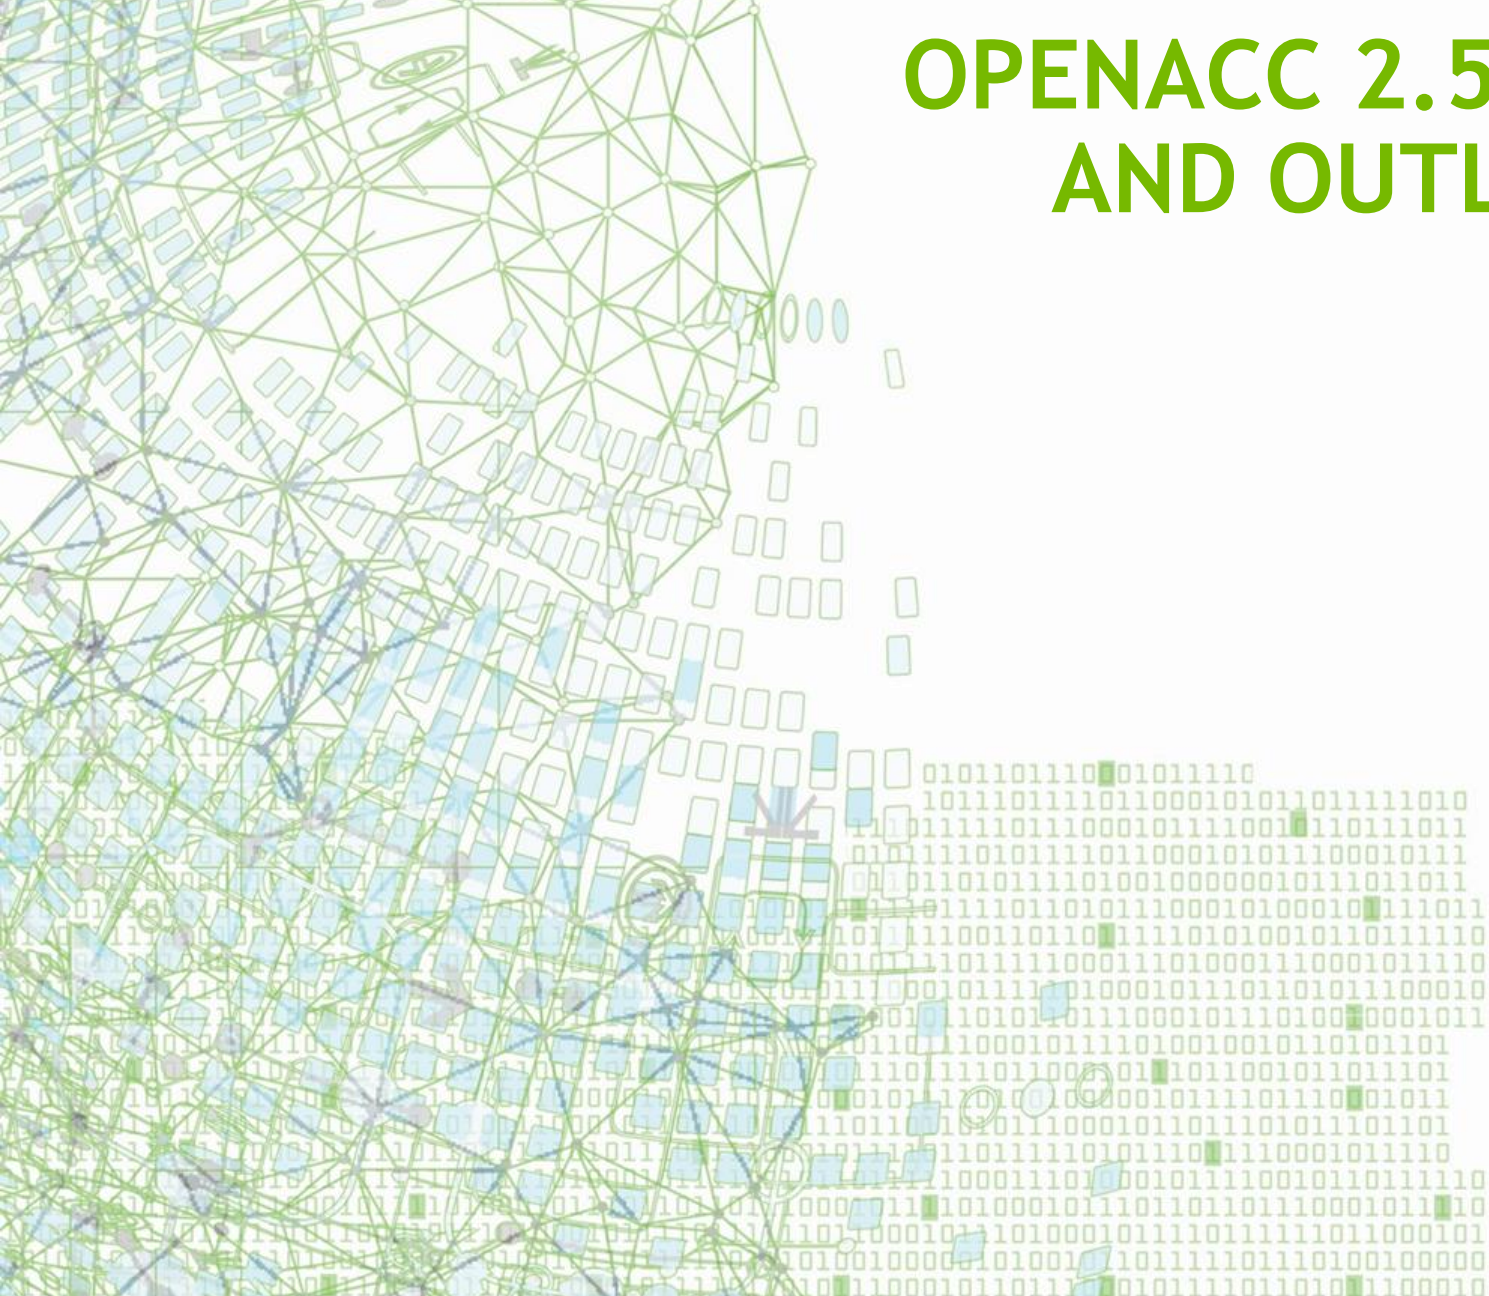
CUDA UM performance is within  
10% -15% of direct-based data  
movement

356.sp anomaly: Fortran  
automatics in a routine with  
accelerated loops

351.palm: lots of paging  
between host and device

PGI 17.1 Compilers OpenACC SPEC ACCEL™ 1.1 performance measured March, 20167 SPEC® and the benchmark name SPEC ACCEL™ are registered trademarks of the Standard Performance Evaluation Corporation.

# OPENACC 2.5 SPECIFICATION AND OUTLOOK ON 2.6





# OPENACC 2.5 FEATURES

(incomplete, see specification for full list)

- `num_gangs`, `num_workers`, `vector_length` are now allowed on `kernels` construct
- **New directives:** `init`, `shutdown`, `set`
- **New `if_present` clause for `update` directive**
- `acc routine` **without** `gang/worker/vector/seq` is now an error
- **New `default(present)` clause was added for `compute` constructs**
- **New API routines `acc_get_default_async`, `acc_set_default_async`**
- Asynchronous versions of data API routines added
- new OpenACC interface for profile and trace tools was added
- various clarifications added to the text

# OPENACC 2.5 FEATURES (CONT.)

(incomplete, see specification for full list)

- Reference counting is now used to manage the correspondence and lifetime of device data
- Data clauses `copy[in/out]`, `create` changed to behave as `present_or_copy[in/out]`, `present_or_create`. Similar change for API routines `acc_copyin`, etc.
- The behavior of the `exit data` directive has changed to decrement the dynamic reference count. A new optional `finalize` clause was added to set the dynamic reference count to zero.

```
#pragma acc data copy(var)
{...
    #pragma acc data copy(var)
    {...}
}
```

```
Reference count for var: 0
1, copyin

2, no copyin!
1 after }, no copyout!
0, copyout
```

# LIKELY OPENACC 2.6 FEATURES

(2.6 specification not finalized yet!)

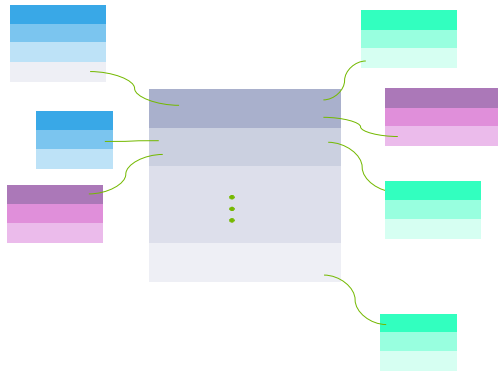
- manual deep copy in the spec (already supported by PGI)
- `no_create` clause (user request)
- `if` and `if_present` clauses on `host_data` directive (user request)
- standardize behavior of Fortran optional arguments in data clauses (user request)
- Fortran bindings for all API routines (user request)
- serial offload region (user request)
- device query routines
- method to catch errors and fail gracefully (user request; this is not error recovery, but would allow an MPI program to shut down)
- a number of small cleanup items

# DEEP COPY

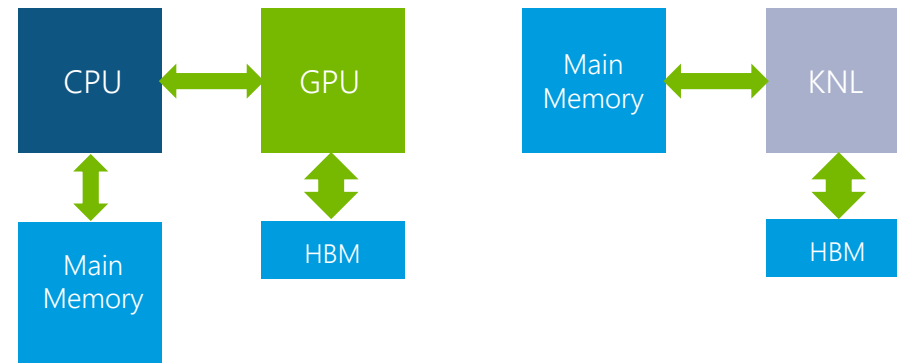


# OPENACC 3.0 DEEP COPY

Modern Data Structures



Modern HPC Node Architectures





# MANUAL DEEP COPY

```
typedef struct points {  
    float* x; float* y; float* z;  
    int n;  
    float coef, direction;  
} points;
```

```
void sub ( int n, float* y ) {  
    points p;  
    #pragma acc data create (p)  
    {  
        p.n = n;  
        p.x = ( float*) malloc ( sizeof ( float )*n );  
        p.y = ( float*) malloc ( sizeof ( float )*n );  
        p.z = ( float*) malloc ( sizeof ( float )*n );  
        #pragma acc update device (p.n)  
        #pragma acc data copyin (p.x[0:n], p.y[0: n])  
        {  
            #pragma acc parallel loop  
            for ( i =0; i<p.n; ++i ) p.x[i] += p.y[i];  
            . . .  
        }  
    }
```

# TRUE DEEP COPY

```
typedef struct points {
    float* x; float* y; float* z;
    int n;
    float coef, direction;
    #pragma acc policy inout(x[0:n],y[0:n])
} points;

void sub ( int n, float* y ) {
    points p;

    p.n = n;
    p.x = ( float*) malloc ( sizeof ( float )*n );
    p.y = ( float*) malloc ( sizeof ( float )*n );
    p.z = ( float*) malloc ( sizeof ( float )*n );

    #pragma acc data copy (p)
    {
        #pragma acc parallel loop
        for ( i =0; i<p.n; ++i ) p.x[i] += p.y[i];
        . . .
    }
}
```

# QUESTIONS ?

Check out [www.openacc.org](http://www.openacc.org) for online resources, news, events, etc.

Join the OpenACC usergroup, e.g. through  
<https://www.openacc.org/community>